

Time Series Modeling

1. Plot the data with proper labeling and make some observations on the graph.

In [65]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from datetime import datetime
from sklearn import metrics
from statsmodels.tsa.ar_model import AutoReg
from numpy import sqrt
from sklearn.metrics import mean_squared_error
```

In ...

```
time_df = pd.read_csv("/Users/theranmeadows/Desktop/Bellevue University/dsc630pred
time_df.head()
```


Out[...]

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP
0	1992	146925	147223	146805	148032	149010	149800	150761.0	151067.0	152588.0
1	1993	157555	156266	154752	158979	160605	160127	162816.0	162506.0	163258.0
2	1994	167518	169649	172766	173106	172329	174241	174781.0	177295.0	178787.0
3	1995	182413	179488	181013	181686	183536	186081	185431.0	186806.0	187366.0
4	1996	189135	192266	194029	194744	196205	196136	196187.0	196218.0	198859.0

In [46]:

```
time_df.plot('YEAR', 'DEC')
```

Out[46]:<Axes: xlabel='YEAR'>



Although not pictured to save room, plotting all months by year showed a very similar line shape across all months. An increase from 1995 to 2006-2007, a sharp decline from 2008-2009 and then back to an increase all the way till 2021. We can confirm this with the MELT function below.

```
In [4...
#transpose the data with melt
time_df_short = pd.melt(time_df, id_vars=['YEAR'], var_name="MONTH", value_name='
```

```
In...
#assign number to the month names
time_df_short['Month'] = time_df_short['MONTH'].map({'JAN': 1, 'FEB': 2, 'MAR': 3,
```

```
In [4...
time_df_short['DATE']=pd.to_datetime(time_df_short[['YEAR', 'Month']].assign(DAY=
```

```
In [50]:
time_df_short.head()
```

```
Out[50]:
```

	YEAR	MONTH	SALES	Month	DATE
0	1992	JAN	146925.0	1	1992-01-01
1	1993	JAN	157555.0	1	1993-01-01
2	1994	JAN	167518.0	1	1994-01-01
3	1995	JAN	182413.0	1	1995-01-01
4	1996	JAN	189135.0	1	1996-01-01

```
In [51]:
#clean the data up
time_df_clean = time_df_short[['DATE', 'SALES']].sort_values('DATE')
time_df_clean = time_df_clean.dropna()
```

```
time_datetime_index = pd.DatetimeIndex(time_df_clean['DATE'].values)
time_df_clean = time_df_clean.set_index(time_datetime_index)
time_df_clean.drop('DATE', axis=1, inplace=True)
time_df_clean
```


Out[51]:

	SALES
1992-01-01	146925.0
1992-02-01	147223.0
1992-03-01	146805.0
1992-04-01	148032.0
1992-05-01	149010.0
...	...
2021-02-01	504458.0
2021-03-01	559871.0
2021-04-01	562269.0
2021-05-01	548987.0
2021-06-01	550782.0

354 rows × 1 columns

In [52]:

```
#plot the new frame to confrim the line trends of the first graph
plt.plot(time_df_clean["SALES"])
plt.title('Sales by Month')
plt.xlabel('Date')
plt.ylabel('Sales in $')
plt.show()
```



2. Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set.

```
In [53]: train_time_df = time_df_clean[time_df_clean.index < '2020-07-01']  
         test_time_df = time_df_clean[time_df_clean.index >= '2020-07-01']
```

```
In [54]: train_time_df
```

```
Out[54]:
```

	SALES
1992-01-01	146925.0
1992-02-01	147223.0
1992-03-01	146805.0
1992-04-01	148032.0
1992-05-01	149010.0
...	...
2020-02-01	459610.0
2020-03-01	434281.0
2020-04-01	379892.0
2020-05-01	444631.0

SALES**2020-06-01** 476343.0

342 rows × 1 columns

In [55]:
test_time_dfOut[55]: **SALES****2020-07-01** 481627.0**2020-08-01** 483716.0**2020-09-01** 493327.0**2020-10-01** 493991.0**2020-11-01** 488652.0**2020-12-01** 484782.0**2021-01-01** 520162.0**2021-02-01** 504458.0**2021-03-01** 559871.0**2021-04-01** 562269.0**2021-05-01** 548987.0**2021-06-01** 550782.0

3. Use the training set to build a predictive model for the monthly retail sales.

In [58]:
time_model = AutoReg(train_time_df, lags=5).fit()

/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

self._init_dates(dates, freq)

In [59]:
print(time_model.summary())**AutoReg Model Results**

```
=====
Dep. Variable:          SALES    No. Observations:          342
Model:                AutoReg(5)  Log Likelihood          -3375.942
Method:              Conditional MLE  S.D. of innovations      5424.567
Date:                Sat, 03 Feb 2024  AIC                        6765.884
Time:                  20:36:22    BIC                      6792.624
Sample:              06-01-1992    HQIC                     6776.542
                  - 06-01-2020
=====
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const      2180.0628    1089.933      2.000      0.045     43.833     4316.293
=====
```

SALES.L1	0.8523	0.055	15.449	0.000	0.744	0.960
SALES.L2	-0.3524	0.080	-4.407	0.000	-0.509	-0.196
SALES.L3	0.3427	0.130	2.629	0.009	0.087	0.598
SALES.L4	0.2459	0.134	1.839	0.066	-0.016	0.508
SALES.L5	-0.0908	0.100	-0.904	0.366	-0.288	0.106


Roots

	Real	Imaginary	Modulus	Frequency
AR.1	1.0013	-0.0000j	1.0013	-0.0000
AR.2	0.0573	-1.2351j	1.2364	-0.2426
AR.3	0.0573	+1.2351j	1.2364	0.2426
AR.4	-2.0016	-0.0000j	2.0016	-0.5000
AR.5	3.5941	-0.0000j	3.5941	-0.0000

4. Use the model to predict the monthly retail sales on the last year of data.

```
In [... time_pred = time_model.predict(start = len(train_time_df), end = (len(time_df_cle
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/deterministic.py:302: UserWarning: Only PeriodIndexes, DatetimeIndexes with a frequency set, RangesIndexes, and Index with a unit increment support extending. The index is set will contain the position relative to the data length.
    fcast_index = self._extend_index(index, steps, forecast_index)
In [61]: plt.plot(time_pred)
plt.plot(test_time_df)
```

Out[61]:[<matplotlib.lines.Line2D at 0x1200c6690>]



5. Report the RMSE of the model predictions on the test set.

```
In [67]: rmse = sqrt(mean_squared_error(test_time_df, time_pred))  
         print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 73875.521